

8 Capabilities

8.1 Evaluation summary

Claude Fable 5.1 is more capable than Fable 5 and achieves state of the art performance on many benchmarks. All the evaluations in this section were run on the final snapshot of Claude Fable 5.1 or Claude Mythos 5.1.

Evaluation		Fable family models		Other models	
		Claude Fable 5.1/ Mythos 5.1	Claude Fable 5/ Mythos 5	Claude Opus 5	GPT-5.6 Sol
SWE-bench Pro		81.2	80	79.2	64.6
SWE-bench Multilingual		89.1	86.6	89.5	-
SWE-bench Multimodal		54.7	54.1	59.4	-
Terminal-Bench 4.0		56% (61%)	42% (45%)	52%	37%
Terminal-Bench-Science 0.1		52.6%	24.7%	29.0%	22.4%
Humanity's Last Exam	No tools	60.9%	57.8%	56.6%	-
	With tools	65.0%	63.8%	63.6%	-
OSWorld 2.0 (partial/strict)		77.9/41.7	72.9/36.1	75.4/39.6	-
HealthBench Professional		62.1%	63.3%	59.8%	-
GDPval-AA v2		1853	1723	1824	1711
AA-Briefcase		1694	1572	1685	1502
AutomationBench		31.4	17.1	26.9	19.6
ARC-AGI-1		97.5%	98.5%	97.5%	96.5%
ARC-AGI-2		90.0%	89.2%	90.42%	92.5%

[Table 8.1.A] **Capability evaluation summary.** Unless otherwise noted, all results for Claude Fable 5.1 and Claude Mythos 5.1 use the following standard configuration: adaptive thinking at max effort, default sampling settings

(temperature, top_p), averaged over five trials. Context window sizes are evaluation dependent and do not exceed 1M tokens. The best score in each row is **bolded**. Competitor figures are drawn from the respective developers' published system cards or benchmark leaderboards. See the [Claude Opus 5 System Card](#) for evaluation details of earlier Claude models.

8.2 SWE-bench Pro, Multilingual, and Multimodal

SWE-bench tests AI models on real-world software engineering tasks. We report three variants, where each score is an average over five trials:

- SWE-bench **Pro**¹² is composed of problems drawn from actively maintained repositories with large, multi-file diffs. Fable 5.1 achieved 81.2%.
- SWE-bench **Multilingual** extends the format to 300 problems across nine programming languages. Fable 5.1 achieved 89.1%.
- SWE-bench **Multimodal**¹³ adds visual context (screenshots, design mockups) to the issue descriptions. Fable 5.1 achieved 54.7%.

8.3 DeepSWE v1.1

DeepSWE is a set of 113 long-horizon software engineering tasks built to measure the capabilities of frontier coding agents. The tasks are diverse, reflect real-world complexity, and are written from scratch to avoid benchmark contamination. Fable 5.1 scored an average of 67.4% over five trials.

A note on scoring: DeepSWE grades each task with hidden tests that are often written for a single reference solution. When we reviewed failing transcripts, we found that Fable 5.1 implemented some ambiguous tasks more thoroughly than the task required, for example by validating inputs and raising exceptions, or by keeping new code consistent with codebase conventions. This resulted in equally valid or more rigorous implementations failing the hidden tests.

8.4 FrontierCode

FrontierCode is an agentic coding benchmark of 150 software engineering tasks created by Cognition. Tasks are derived from real pull requests in open-source repositories, such as fixing websocket bugs in [aiohttp](#), hardening Prisma's browser bundle, or extending JSON schema linting rules. Each task gives the agent a checked-out repository and a single issue

¹² Deng, X., et al. (2025). SWE-Bench Pro: Can AI agents solve long-horizon software engineering tasks? arXiv:2509.16941. <https://arxiv.org/abs/2509.16941>

¹³ Yang, J., et al. (2024). SWE-bench Multimodal: Do AI Systems generalize to visual software domains? arXiv:2410.03859. <https://arxiv.org/abs/2410.03859>

results on some tasks, but drops off more often on mid-difficulty tasks, whereas Fable 5.1's performance holds up as task difficulty increases. Fable 5.1 also has a substantially higher floor than Fable 5, with a median task score of 0.56 versus 0.41, and the lowest outright failure rate of the three (5% of trials, versus 6% for Opus 5 and 8% for Fable 5). It has the highest share of top-end results as well, with 38% of trials scoring above 0.8. Fable 5.1 leads on breadth and consistency: it combines high-end capability with a stronger performance floor, particularly on difficult, long-horizon engineering work.

8.6 Terminal-Bench 4.0

Terminal-Bench 4.0 tests AI models on real-world work in terminal and command-line containerized environments. It's a set of 66 tasks, with an emphasis on science-adjacent and frontier engineering problems like computational biology, physics simulation, CAD, formal proofs, and GPU performance work. Previous Terminal-Bench versions were [very timeout- and resource-sensitive](#); 4.0 largely addresses these issues by increasing timeouts and adaptively bumping up both RAM and CPU in a few tasks that could benefit from it. The choice of sandboxing provider might still be a confounding factor in a few tasks, but we expect it to play a significantly smaller role than it did in previous versions. We also found that the upgrade has reduced the confounding role of various harnesses (e.g., CLI memory footprint, compaction strategy, and container communication protocol), ultimately making Terminal-Bench a more stable and rigorous evaluation.

Claude Mythos 5.1 scored 60.9% on Terminal-Bench 4.0, while Claude Fable 5.1 scored 55.8%. Claude Opus 5 scored 52.3%, and Claude Fable 5 42.0%. Scores are averaged over 10 trials per task (660 trials) for Mythos 5.1 and 15 trials per task (990 trials) for the other three models; the standard error is ± 1.6 -2 points for every model. The numbers reported use Claude Code in --bare mode and maximum thinking effort. For reference, the public leaderboard reports Opus 5 at 51.8% and Fable 5 at 44.5% (five trials, Claude Code harness), within the noise of our internal setup. The public leaderboard also reports OpenAI's GPT-5.6 Sol at 37.3%, scored with the Codex CLI at max thinking effort.

8.7 Terminal-Bench-Science 0.1

Terminal-Bench-Science is a Stanford-led community benchmark of 70 tasks drawn from scientific research workflows, authored and reviewed by scientists and researchers across the life, physical, earth, mathematical, and engineering sciences. Scientific advisors to the project are from MIT, Princeton, University of Washington, Genentech, and Stanford.