

8 Capabilities

8.1 Evaluation summary

Evaluation		Claude family models				Other models	
		Mythos 5	Fable 5	Mythos Preview	Opus 4.8	GPT-5.5	Gemini 3.1 Pro
SWE-bench Pro		80.3	80	77.8	69.2	58.6	54.2
SWE-bench Verified		95.5	95	93.9	88.6	-	80.6
Terminal-Bench 2.1		88.0	84.3	-	82.7	83.4 (Codex CLI)	70.7 (Gemini CLI)
BrowseComp		88.0 (single-agent) 93.3 (multi-agent)	-	87.9	84.3 (single-agent) 88.5 (multi-agent)	84.4	85.9
Humanity's Last Exam	No tools	59.0	-	56.8	49.8	41.4	44.4
	With tools	64.5	-	64.7	57.9	52.2	51.4
CharXiv Reasoning	No tools	88.9	-	86.2	80.5	-	-
	With tools	93.5	-	92.5	89.9	-	-
BioMystery Bench	Human	83.9	-	82.6	80.4		
	Hard	46.1	-	29.6	40		
OSWorld-Verified ²⁷		85.0	85.0	85.4	83.4	78.7	76.2 (3.5 Flash: 78.4)
CritPt		28.6	-		20.9	27.1	17.7
ArxivMath		78.5		68.7	71.8	71.5	64.8

²⁷ Changes to the Mythos OSWorld score are due to a bug fix on our zoom tool when paired with batched actions, and increasing the max tokens per turn from 16K to 128K.

8.2 SWE-bench Verified, Pro, Multilingual, and Multimodal

SWE-bench (Software Engineering Bench) tests AI models on real-world software engineering tasks. We report four variants, where the score is the average over 5 trials:

- SWE-bench **Verified**²⁹ is a 500-problem subset, each verified by human engineers as solvable. Mythos 5 achieved 95.5% and Fable 5 achieved 95%.
- SWE-bench **Pro**³⁰ is a harder variant: problems drawn from actively-maintained repositories with larger, multi-file diffs and reduced public ground-truth leakage. Claude Mythos 5 achieved 80.3% and Claude Fable 5 achieved 80%.
- SWE-bench **Multilingual** extends the format to 300 problems across 9 programming languages. Mythos 5 achieved 92.2%.
- SWE-bench **Multimodal**³¹ adds visual context (screenshots, design mockups) to the issue descriptions (see Section 9.3 of the [Claude Opus 4.7 System Card](#) for details on the internal harness). Mythos 5 achieved 54.9%.


All SWE-bench variants use the standard configuration, with thinking blocks included in the sampling results. For our memorization screening, see Section 6.2.1 in the [Mythos Preview System Card](#).

²⁹ Jimenez, C. E., et al. (2024). SWE-bench: Can language models resolve real-world GitHub issues? arXiv:2310.06770. <https://arxiv.org/abs/2310.06770>

³⁰ Deng, X., et al. (2025). SWE-Bench Pro: Can AI agents solve long-horizon software engineering tasks? arXiv:2509.16941. <https://arxiv.org/abs/2509.16941>

³¹ Yang, J., et al. (2024). SWE-bench Multimodal: Do AI systems generalize to visual software domains? arXiv:2410.03859. <https://arxiv.org/abs/2410.03859>

SWE-bench Pro: test-time compute scaling



[Figure 8.2.A] SWE-bench Pro score versus average cost per task across reasoning-effort levels.

8.3 Terminal-Bench 2.1

Terminal-Bench 2.1³² tests AI models on real-world coding tasks in terminal and command-line environments. We've decided to switch to a new harness, [mini-SWE-agent](#), which is more robust to timeouts compared to the Terminus-2 harness that we've previously reported. For example, at xhigh effort, Terminus-2 experiences 2.7× more timeouts than mini-SWE-agent, due to the way it waits for commands execution through a tmux session; this makes final scores noisier and less legible.

Using the mini-SWE harness, with a [GKE cluster](#), 1× timeout rate and 3× memory ceiling before pod preemption:

- Claude Mythos 5: achieved 88% mean reward, averaged over 5 attempts for each one of the 89 unique tasks (for a total of 445 trials), at high effort.
- Claude Fable 5: achieved 84.3% mean reward—with 20.9% of trials hitting a safety refusal and falling back to Claude Opus 4.8 for the rest of the trajectory, at high effort.

³² Merrill, M. A., et al. (2026). Terminal-Bench: Benchmarking agents on hard, realistic tasks in command line interfaces. arXiv:2601.11868. <https://arxiv.org/abs/2601.11868>