

6.2.3 MMMU-Pro

MMMU-Pro is a benchmark we would normally report in this System Card (specifically, in Section 6.11 below). Like CharXiv Reasoning, MMMU-Pro comprises material from widely disseminated public materials—for example, university exams, textbooks, and quiz sites—that are difficult to fully decontaminate from training corpora.

We identified a large fraction of MMMU-Pro images that appear in the training data, primarily via textbooks, homework-help sites, and document crawls, which repackage and distribute the underlying source content.

Unlike CharXiv Reasoning, MMMU-Pro contains a limited number of questions for which variants of *approximately equivalent difficulty* can be readily created. MMMU-Pro contains a small number of charts and figures, but studying just this subset of problems would paint a biased picture. Given the difficulty of determining the impact of contamination, we choose to omit results for MMMU-Pro from this System Card.

6.3 Overall results summary

Table 6.3.A summarizes the evaluations discussed in more detail below.

Evaluation	Claude family		Other models	
	Claude Mythos Preview	Claude Opus 4.6	GPT-5.4	Gemini 3.1 Pro
SWE-bench Verified	93.9%	80.8%	—	80.6%
SWE-bench Pro	77.8%	53.4%	57.7%	54.2%
SWE-bench Multilingual	87.3%	77.8%	—	—
SWE-bench Multimodal	59%	27.1%	—	—
Terminal-Bench 2.0*	82%	65.4%	75.1%	68.5%
GPQA Diamond	94.5%	91.3%	92.8%	94.3%

Evaluation		Claude family		Other models	
		Claude Mythos Preview	Claude Opus 4.6	GPT-5.4	Gemini 3.1 Pro
MMMLU		92.7%	91.1%	—	92.6%–93.6%
USAMO 2026		97.6%	66.2%	95.2%	74.4%
GraphWalks BFS 256K-1M		80.0%	38.7%	21.4%	—
HLE	no tools	56.8%	40.0%	39.8%	44.4%
	with tools	64.7%	53.1%	52.1%	51.4%
CharXiv Reasoning	no tools	86.1%	69.1%	–	–
	with tools	93.2%	84.7%	–	–
OSWorld		79.6%	72.7%	75.0%	

[Table 6.3.A] **Capability Evaluation Summary.** Unless otherwise noted, all Claude Mythos Preview results use the following standard configuration: adaptive thinking at max effort, default sampling settings (temperature, top_p), averaged over 5 trials. Context window sizes are evaluation-dependent and do not exceed 1M tokens. The best score in each row is **bolded**. Competitor figures are drawn from the respective developers’ published system cards or benchmark leaderboards. See the [Claude Opus 4.6 System Card](#) for evaluation details of earlier Claude models. *For Terminal-Bench 2.0, OpenAI used a specialized harness for their reported score, making comparison between the models in this row inexact. All other scores used the Terminus-2 harness.

6.4 SWE-bench Verified, Pro, Multilingual, and Multimodal

SWE-bench (Software Engineering Bench) tests AI models on real-world software engineering tasks. We report four variants:

- [SWE-bench Verified](#) (OpenAI) is a 500-problem subset, each verified by human engineers as solvable. Claude Mythos Preview achieves 93.9%, averaged over 5 trials.
- [SWE-bench Pro](#) (Scale) is a harder variant: problems drawn from actively-maintained repositories with larger, multi-file diffs and no public ground-truth leakage. Claude Mythos Preview achieves 77.8%, averaged over 5 trials.

- [SWE-bench Multilingual](#) extends the format to 300 problems across 9 programming languages. Claude Mythos Preview achieves 87.3%, averaged over 5 trials.
- [SWE-bench Multimodal](#) adds visual context (screenshots, design mockups) to the issue descriptions. Claude Mythos Preview achieves 59.0% (evaluated on an internal harness; see Appendix 8.6), averaged over 5 trials. We note higher trial-to-trial variance on this variant (56.4%–61.4%) than on the others.

All SWE-bench variants use the standard configuration (see Table 6.3.A), with thinking blocks included in the sampling results. For our memorization screening, see Section 6.2.

6.5 Terminal-Bench 2.0

[Terminal-Bench 2.0](#), developed by researchers at Stanford University and the Laude Institute, tests AI models on real-world tasks in terminal and command-line environments.

We ran Terminal-Bench 2.0 in the Harbor scaffold with the Terminus-2 harness and default parser. Each task runs in an isolated Kubernetes pod with guaranteed resources at 1× the benchmark-specified limits (hard preemption ceiling at 3×) and timeouts at 1× for benchmark fidelity. Details on this configuration are available at [our engineering blog](#).

Claude Mythos Preview achieved **82% mean reward**, averaged over 5 attempts for each one of the 89 unique tasks (for a total of 445 trials). We configured Claude Mythos Preview to run with maximum reasoning effort (adaptive mode), 1M total token budget per task, and 32K maximum output tokens per request. Terminal-Bench is sensitive to inference latency: fixed wall-clock timeouts mean a slower-decoding endpoint completes fewer episodes per task. Our reported score uses a production API endpoint to account for these dynamics.

Terminal-Bench 2.0 timeouts get quite restrictive at times, especially with thinking models, which risks hiding real capabilities jumps behind seemingly uncorrelated confounders like sampling speed. Moreover, some Terminal-Bench 2.0 tasks have ambiguities and limited resource specs that don't properly allow agents to explore the full solution space—both being currently addressed by the maintainers in the [2.1 update](#). To exclusively measure agentic coding capabilities net of the confounders, we also ran Terminal-Bench with the latest 2.1 fixes available on GitHub, while increasing the timeout limits to 4 hours (roughly four times the 2.0 baseline). This brought the mean reward to **92.1%**. Under the same conditions, we measured GPT-5.4 with Codex CLI harness to achieve 75.3% (up from 68.3% under baseline specs)²³.

²³ We do not report a Gemini 3.1 Pro result with this setup. We struggled to reproduce previous best results, including our own tests, which matched the reported scores when the model was released in February.