

8 Capabilities

8.1 Evaluation summary

Evaluation		Claude family models		Other models	
		Claude Opus 4.8	Claude Opus 4.7	GPT-5.5	Gemini 3.1 Pro
SWE-bench Verified		88.6	87.6	-	80.6
SWE-bench Pro		69.2	64.3	58.6	54.2
SWE-bench Multilingual		84.4	80.5	-	-
SWE-bench Multimodal		38.4	34.5	-	-
BrowseComp ²⁷		84.3 (single-agent) 88.5 (multi-agent)	79.8	84.4	85.9
Terminal-Bench 2.1		74.6	66.1	78.2	70.3
Humanity's Last Exam	No tools	49.8	46.9	41.4	44.4
	With tools	57.9	54.7	52.2	51.4
ChartQAPro	No tools	69.4	67.6	-	-
	With tools	72.3	69.8	-	-
OSWorld-Verified ²⁸		83.4	82.8	78.7	76.2 (3.5 Flash: 78.4)
GPQA Diamond		93.6	94.2	-	94.3
ScreenSpot-Pro	No tools	82.3	79.5	-	-

²⁷ Changes to the Opus 4.7 BrowseComp score are due to a new blocklist, context compaction at 200k tokens and using adaptive thinking.

²⁸ Changes to the Opus 4.7 OSWorld score are due to a bug fix on our zoom tool when paired with batched actions, and increasing the max tokens per turn from 16K to 128K.

ScreenSpot-Pro	With tools	87.9	87.6	-	-
Finance Agent v2		53.9	51.5	51.8	43.0 (3.5 Flash: 57.9)
GDPval-AA		1890	1753	1769	1314
MCP-Atlas		82.2	79.1	75.3	78.2 (Gemini 3.5 Flash: 83.6)
Automation Bench		15.5	9.9	12.9	9.6 (3.5 Flash: 14.5)
GraphWalks BFS 256K		85.9	76.9	73.7	-
GraphWalks Parents 256K		99.3	93.6	90.1	-

[Table 8.1.A] **Capability evaluation summary.** Unless otherwise noted, all Claude Opus 4.8 results use the following standard configuration: adaptive thinking at max effort, default sampling settings (temperature, top_p), averaged over 5 trials. Context window sizes are evaluation-dependent and do not exceed 1M tokens. The best score in each row is **bolded**. Competitor figures are drawn from the respective developers' published system cards or benchmark leaderboards. See the [Claude Opus 4.7 System Card](#) for evaluation details of earlier Claude models.

8.2 SWE-bench Verified, Pro, Multilingual, and Multimodal

SWE-bench (Software Engineering Bench) tests AI models on real-world software engineering tasks. We report four variants, where the score is the average over 5 trials:


- SWE-bench **Verified**²⁹ is a 500-problem subset, each verified by human engineers as solvable. Opus 4.8 achieves 88.6%.
- SWE-bench **Pro**³⁰ is a harder variant: problems drawn from actively-maintained repositories with larger, multi-file diffs and no public ground-truth leakage. Opus 4.8 achieves 69.2%.
- SWE-bench **Multilingual** extends the format to 300 problems across 9 programming languages. Opus 4.8 achieves 84.4%.

²⁹ Jimenez, C. E., et al. (2024). SWE-bench: Can Language Models Resolve Real-World GitHub Issues? arXiv:2310.06770. <https://arxiv.org/abs/2310.06770>

³⁰ Deng, X., et al. (2025). SWE-Bench Pro: Can AI Agents Solve Long-Horizon Software Engineering Tasks? arXiv:2509.16941. <https://arxiv.org/abs/2509.16941>

- SWE-bench **Multimodal**³¹ adds visual context (screenshots, design mockups) to the issue descriptions. Opus 4.8 achieves 38.4% (evaluated on an internal harness; Section 9.3 of the [Claude Opus 4.7 System Card](#)).

All SWE-bench variants use the standard configuration, with thinking blocks included in the sampling results. For our memorization screening, see Section 6.2.1 in the [Mythos Preview System Card](#).



[Figure 8.2.A] SWE-Bench Pro pass rate across reasoning effort levels with mean output tokens per task on a log scale. Opus 4.8 reaches its peak pass rate at extra-high effort, with maximum effort performing comparably. At minimum effort, Opus 4.8 matches the peak performance of Opus 4.7 at maximum effort.

8.3 Terminal-Bench 2.1

Terminal-Bench 2.1³² tests AI models on real-world tasks in terminal and command-line environments. The third-party evaluation framework Harbor runs it on the Daytona cloud provider, using the official integration. On Harbor's leaderboard, using the Terminus-2 harness, Opus 4.8 achieved 74.6% mean reward, averaged over 5 attempts for each one of the 89 unique tasks (for a total of 445 trials) compared to Opus 4.7's score of 66.1%. We

³¹ Yang, J., et al. (2024). SWE-bench Multimodal: Do AI Systems Generalize to Visual Software Domains? arXiv:2410.03859. <https://arxiv.org/abs/2410.03859>

³² Merrill, M. A., et al. (2026). Terminal-Bench: Benchmarking Agents on Hard, Realistic Tasks in Command Line Interfaces. arXiv:2601.11868. <https://arxiv.org/abs/2601.11868>