

Evaluation		Claude family models				Other models	
		Claude Sonnet 4.6	Claude Opus 4.6	Claude Opus 4.5	Claude Sonnet 4.5	Gemini 3 Pro	GPT-5.2 (all models)
SWE-bench Verified ²		79.6%	80.8%	80.9%	77.2%	76.2%	80.0%
Terminal-Bench 2.0		59.1% (default thinking)	65.4%	59.8%	51.0%	56.2%	64.7%
τ^2 -bench	Retail	91.7%	91.9%	88.9%	86.2%	85.3%	82.0%
	Telecom	97.9%	99.3%	98.2%	98.0%	98.0%	98.7%
MCP-Atlas		61.3%	59.5% ³	62.3%	43.8%	54.1%	60.6%
OSWorld-Verified		72.5%	72.7%	66.3%	61.4%	—	—
ARC-AGI-2 (Verified)		58.3% ⁴	68.8% ³	37.6%	13.6%	31.1%	54.2%
GPQA Diamond		89.9%	91.3%	87.0%	83.4%	91.9%	93.2%
MMMLU		89.3%	91.1%	90.8%	89.5%	91.8%	89.6%
GDPval-AA		1633	1606	1416	1276	1201	1462
MMMU-Pro	No tools	74.5%	73.9%	70.6%	63.4%	81%	79.5%
	With tools	75.6%	77.3%	73.9%	68.9%	—	80.4%
HLE	No tools	33.2%	40.0%	30.8%	17.7%	37.5%	36.6%
	With tools	49.0%	53.0%	43.4%	33.6%	45.8%	50.0%

[Table 2.1.A] All Claude Sonnet 4.6 evaluation results are an average over 10 trials unless otherwise noted. Each run uses adaptive thinking, max effort, and default sampling settings (temperature, top_p). Context window sizes are evaluation-dependent, but do not ever exceed 1M. The best score in each row is **bolded**.

² SWE-bench results are averaged over 25 trials.

³ We report the max effort score in this table; with high effort, Claude Opus 4.6 achieves a score of 62.7% (on MCP-Atlas) and 69.2% (on ARC-AGI-2).

⁴ We report the max effort score in this table; with high effort, Claude Sonnet 4.6 achieves a score of 60.42%.

2.2 SWE-bench (Verified and Multilingual)

SWE-bench (Software Engineering Bench) tests AI models on real-world software engineering tasks.

For the [SWE-bench Verified](#) variant, developed by OpenAI, models are shown 500 problems that have been verified by human engineers to be solvable. We also assessed the model on [SWE-bench Multilingual](#). Here, “multilingual” refers to different programming languages: this variant assesses models on their solutions to 300 problems in 9 different languages.

Claude Sonnet 4.6 achieved 79.6% on SWE-bench Verified and 75.9% on SWE-bench Multilingual. Our SWE-bench results are averaged over 10 trials, each run with adaptive thinking, max effort, default sampling settings (temperature, top_p), and with the thinking blocks included in the sampling results.

For SWE-bench Verified, we found that the following prompt modification resulted in a score of 80.2%:

You should use tools as much as possible, ideally more than 100 times. You should also implement your own tests first before attempting the problem. You should take time to explore the codebase and understand the root cause of issues, rather than just fixing surface symptoms. You should be thorough in your reasoning and cover all edge cases.

2.3 Terminal-Bench 2.0

[Terminal-Bench 2.0](#), developed by researchers at Stanford University and the Laude Institute, tests AI models on real-world tasks within terminal or command-line environments.

We ran Terminal-Bench 2.0 in the Harbor scaffold using the Terminus-2 harness with the default parser. All experiments described below, including those testing non-Claude models, ran on a GKE cluster using n2-standard-32 nodes (32 vCPUs, 128 GB RAM, 500 GB persistent disk), in us-central1. Each task runs in an isolated Kubernetes pod; guaranteed resource allocation is set at 1× the benchmark-specified limits, with a hard preemption ceiling at 3×. Timeouts are kept at 1×. Details on this configuration and the rationale behind it are available on [our Engineering blog](#).

Using this setup, Claude Sonnet 4.6 achieved a 59.1% pass rate with no thinking budget, and without setting effort level (i.e. max effort). We ran all 89 tasks 5 times each.