

8 Capabilities

8.1 Evaluation summary

Claude Opus 5 is meaningfully more intelligent than Opus 4.8 and achieves state of the art performance on many benchmarks.

Evaluation		Claude models			Other models
		Opus 5	Opus 4.8	Fable 5	GPT 5.6 Sol
SWE-bench Pro		79.2	69.2	80	64.6
SWE-bench Multilingual		89.5	84.4	86.6	-
SWE-bench Multimodal		59.4	38.4	54.1	-
DeepSWE v1.1		68.8	59.0	69.7	72.7
FrontierCode 1.1 (Main)		53.4	46.5	53.5	47.5
FrontierBench v0.1		43.3	21.1	33.8	34.4 (Codex)
BrowseComp		90.8	84.3	87.4	90.4
Humanity's Last Exam	No tools	56.3	49.8	56.5	-
	With tools	64.7	57.9	63.9	-
OSWorld 2.0		70.6	55.7	66.1	62.6
HealthBench Professional		59.8	57.4	66.0¹¹	60.5
GDPval-AA v2		1861	1593	1747	1736
AA-Briefcase		1720	1346	1574	1505
AutomationBench		26.0	17.0	17.4	18.1
ARC-AGI-1		97.5	92.5	-	97.5 (xhigh)

¹¹ Mythos 5

ARC-AGI-2	90.4	72.1	-	92.5
ARC-AGI-3	30.2 (high)	1.5	-	7.8

[Table 8.1.A] **Capability evaluation summary.** Unless otherwise noted, all Claude Opus 5 results use the following standard configuration: adaptive thinking at max effort, default sampling settings (temperature, top_p), averaged over 5 trials. Context window sizes are evaluation-dependent and do not exceed 1M tokens. The best score in each row is **bolded**. Competitor figures are drawn from the respective developers' published system cards or benchmark leaderboards. See the [Claude Table 5 System Card](#) for evaluation details of earlier Claude models. FrontierBench results in this table are from Harbor's evaluations.

8.2 SWE-bench Verified, Pro, Multilingual, and Multimodal

SWE-bench (Software Engineering Bench) tests AI models on real-world software engineering tasks. We report four variants, where each score is an average over five trials:

- SWE-bench **Verified**¹² is a 500-problem subset, each verified by human engineers as solvable. Claude Opus 5 achieved 96.0%.
- SWE-bench **Pro**¹³ is a harder variant composed of problems drawn from actively-maintained repositories with larger, multi-file diffs and reduced public ground-truth leakage. Claude Opus 5 achieved 79.2%.
- SWE-bench **Multilingual** extends the format to 300 problems across 9 programming languages. Claude Opus 5 achieved 89.5%.
- SWE-bench **Multimodal**¹⁴ adds visual context (screenshots, design mockups) to the issue descriptions (see [Section 9.3](#) for details on the internal harness). Claude Opus 5 achieved 59.4%.

8.3 DeepSWE v1.1


DeepSWE is a set of 113 long-horizon software engineering tasks built to measure the capabilities of frontier coding agents. The tasks are diverse, reflect real-world complexity, and are written from scratch to avoid benchmark contamination. Claude Opus 5 scored an average of 68.8% over five trials.

¹² Jimenez, C. E., et al. (2024). SWE-bench: Can language models resolve real-world GitHub issues? arXiv:2310.06770. <https://arxiv.org/abs/2310.06770>

¹³ Deng, X., et al. (2025). SWE-bench Pro: Can AI agents solve long-horizon software engineering tasks? arXiv:2509.16941. <https://arxiv.org/abs/2509.16941>

¹⁴ Yang, J., et al. (2024). SWE-bench Multimodal: Do AI systems generalize to visual software domains? arXiv:2410.03859. <https://arxiv.org/abs/2410.03859>

FrontierCode (extended set): test-time compute scaling



[Figure 8.4.B] FrontierCode (extended set). Scores of Claude Opus 5, Claude Opus 4.8, Claude Fable 5, and GPT-5.6 Sol on the extended set of FrontierCode v1.1, a coding evaluation run and scored by Cognition, at each model's five reasoning-effort settings. Claude Opus 5 reaches its best extended-set score, 63.6, at medium effort.

8.5 FrontierBench v0.1

FrontierBench v0.1 is a successor to Terminal-Bench 2.1 developed by the same team. It's a refreshed set of 74 harder tasks, with a larger emphasis on science and engineering problems like computational biology, physics simulation, CAD, formal proofs, and GPU performance work. It tests AI models on real-world work in terminal and command-line containerized environments. In addition to evaluations run by Harbor presented in Table 8.1.A, we report here our own evaluations on the [mini-SWE-agent](#) harness on a GKE backend.

On FrontierBench v0.1, Claude Opus 5 achieved a 44.4% mean reward, averaged over 5 attempts for each one of the 74 unique tasks, using **xhigh** effort (best result – **max** scored similarly and within noise, landing at 43%). **high** effort achieves 39% mean reward for 19% fewer output tokens on average, while **low** effort achieves 25% mean reward for 64% fewer output tokens on average.

We also ran the benchmark on Claude Fable 5, which achieved 33.7% at **max** effort, Claude Sonnet 5, 17% mean reward, and Claude Opus 4.8, which achieved 18.7%.